

Measuring Bias in COVID-19 News Articles

Andrew Pagtakhan, Kwan Bo Shim, Jazmin Trejo

May 10th, 2021

Contents

1	Preface	1
2	Research Question	2
2.1	Background	2
2.2	Research Question	2
2.3	Use Cases	2
2.4	Methodology	2
2.4.1	Data Sourcing	2
3	Data Cleaning	4
3.1	GNews API	4
3.2	Web Scraping	4
3.2.1	BBC	4
3.2.2	CNN	4
3.2.3	Reuters	4
3.2.4	The Wall Street Journal	4
4	Exploratory Data Analysis	4
4.1	Sentiment Analysis	4
5	Feature Engineering	6
5.1	Average Sentiment and Word counts	6
5.2	Keyword features	6
5.3	Topic features from prior research	7
5.4	Topic Modeling: Latent Dirichlet Allocation	7
6	Measuring Bias Across News Sources	7
6.1	Topic Probabilities	7
6.2	Average Sentiment	9
6.3	Word Count	10
7	Classification	10
8	Results	11
8.1	Conclusion and Next Steps	11
8.2	Limitations:	11

1 Preface

The full code and dataset can be found here: <https://github.com/ajpag/US-News-NLP>

All analysis was completed in R.

2 Research Question

2.1 Background

The means in which information is communicated with regards to the COVID-19 pandemic has had major influence on how we read and learn about the virus through a multitude of media outlets. Some of these major sources include television, YouTube, social media forums, and major news companies. Major news companies in particular carry large influence based on the audiences it can reach. For example, Fox News Channel averaged 2.5 million primetime viewers (8pm - 11pm) in February 2021, and CNN averaged 1.7 million during the same time period. *Source: <https://www.foxnews.com/media/fox-news-finishes-february-most-watched-primetime-network>.*

Since major news companies have a large influence in how information is communicated to its audiences, it is vital to quantify how different these sources are in relation to COVID-19 news. By measuring potential bias in relation to each source, this analysis examines how different major news sources are when reporting on COVID-19. It is also insightful to see if there are underlying patterns such as subtopics within COVID-19 that are published more in news sources over others. This would help identify if there are patterns that are predictive of which news source the article came from.

2.2 Research Question

Are there underlying patterns in news articles related to COVID-19 across major news sources that are suggestive of bias, and are these patterns predictive of which news source it is likely from?

2.3 Use Cases

By quantifying underlying differences on COVID-19 reporting and examining the predictive power of these patterns to identify the news source, this study can be useful for a number of use cases. For example, understanding biases in article text can help the reader understand inherent ideological leanings towards certain news sources, which can help equip them with greater understanding and critical examination of news consumption. From a policy perspective, greater impact and studies could be conducted to place more standardized regulations in an effort to influence more objective reporting. We acknowledge this second use case can be difficult from business and philosophical perspectives, especially in relation to the First Amendment of the US Constitution in relation to the freedom of speech.

2.4 Methodology

2.4.1 Data Sourcing

In order to choose major news sources to analyze, the figure below created by the Pew Research Center shows where on the US political spectrum various news companies fall. In order to get a mix of sources across the conservative and liberal spectrum, while balancing limitations of computing resources, the following news sources were used to procure a dataset.

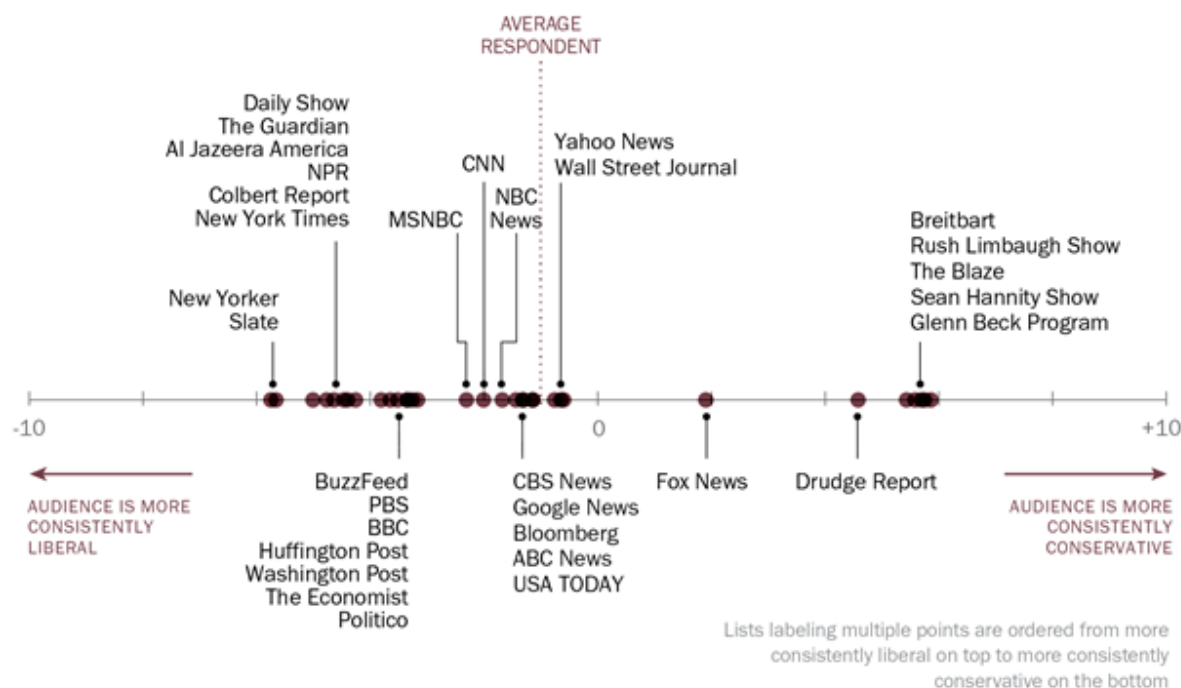
Source: https://www.journalism.org/2014/10/21/political-polarization-media-habits/pj_14-10-21_mediapolarization-08/

- BBC
- CNN
- The Wall Street Journal
- Reuters

Note: Due to legal restrictions, Fox News data was not scraped

Ideological Placement of Each Source's Audience

Average ideological placement on a 10-point scale of ideological consistency of those who got news from each source in the past week...



American Trends Panel (wave 1). Survey conducted March 19-April 29, 2014. Q22. Based on all web respondents. Ideological consistency based on a scale of 10 political values questions (see About the Survey for more details.) ThinkProgress, DailyKos, Mother Jones, and The Ed Schultz Show are not included in this graphic because audience sample sizes are too small to analyze.

PEW RESEARCH CENTER

Figure 1: Source: Pew Research Center

3 Data Cleaning

Leveraging the **GNews API**, news article data related to COVID-19 were pulled for the following major news sources noted in the previous section (filtered to US articles).

3.1 GNews API

A total of 5,360 articles between the time period 1/1/2020 - 4/9/2021 were called from the API (20 articles per week and news source, for 67 weeks). The API call returned 3,454 records for article-related data. Below are some of the key fields from the API call, with an example:

- **article url:** <https://www.cnn.com/2021/02/08/health/covid-19-antigen-tests-states-cnn-analysis/index.html>
- **article description:** “Covid-19 antigen tests not counted among cases in some states, CNN analysis shows - CNN”
- **date and time of article publication:** 2021-02-08T08:00:00Z
- **article source name:** CNN

3.2 Web Scraping

After gathering the article URLs, the full news text for each article was pulled. Each data source carried unique idiosyncrasies in its html structure. The next section highlights unique aspects web scraping each data source.

3.2.1 BBC

3.2.2 CNN

There were two distinct HTML structures. One for the first sentence, and another for the remainder of the article. Additional cleaning needed to be done to remove the “(CNN)” text at the beginning of each article, as well as additional escape characters scattered throughout the article.

3.2.3 Reuters

This was relatively streamlined compared to the other news sources. One unique aspect of Reuters was that a number of their articles were not text articles in the traditional sense, but slideshows. For example this article:

<https://www.reuters.com/news/picture/coronavirus-outbreak-spreads-in-china-idUSRTS2ZART>

3.2.4 The Wall Street Journal

4 Exploratory Data Analysis

4.1 Sentiment Analysis

To better understand the data, sentiment analysis utilizing various lexicons were explored. Key visualizations of the data are illustrated:

BBC had the least average words per article, whereas CNN had the most words per article.

BBC News had a larger variety of top words compared to the other news sources.

All news sources had average weekly sentiment mainly hovering between -1 to 0, using the Afinn lexicon. This means that average sentiment was slightly negative.




Figure 2: Words Per Article




Figure 3: Top Words with a sentiment

The Wall Street Journal had the largest range in average weekly sentiment. It is plausible there is some correlation between this sentiment and the stock market, and would require further analysis outside of this study.

Afinn lexicon: <https://www.tidytextmining.com/sentiment.html>

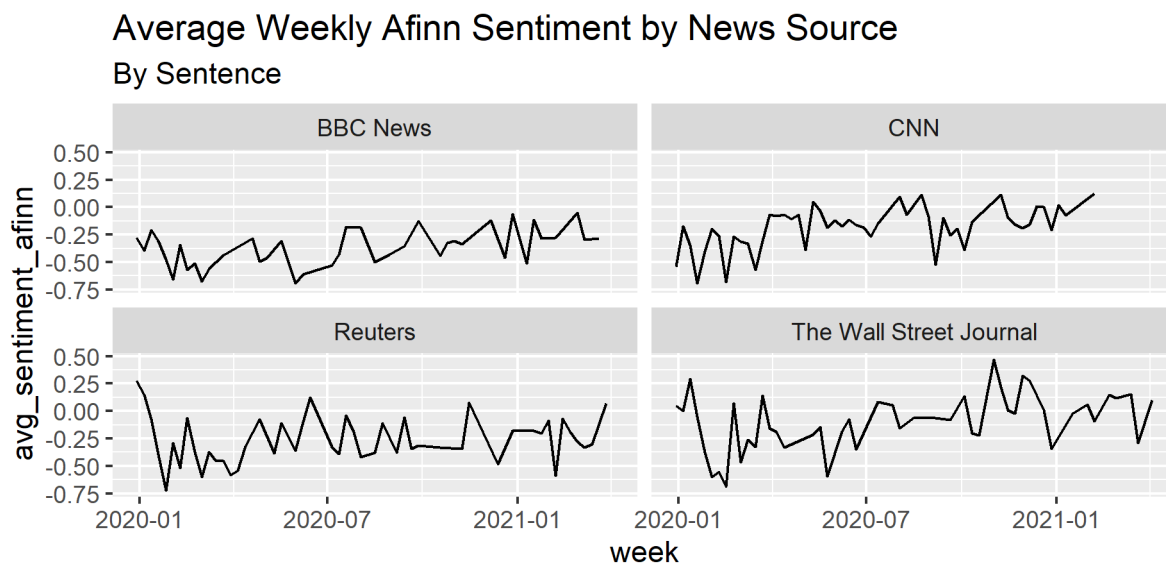


Figure 4: Average Sentiment by Week

5 Feature Engineering

In order to create features to build classification models for predicting a news source based on text, the below features were generated.

5.1 Average Sentiment and Word counts

One feature was generated for each article:

- average sentiment by word
- average sentiment by sentence
- word count
- word count with a sentiment

5.2 Keyword features

As a baseline, keyword features based on the researchers' prior knowledge on the following topics was used to generate average Afinn sentiment for articles. One feature was generated for each term below (noted in quotes).

- Politics: "Trump", "Biden"
- Business: "stock market", "financial"
- Pandemic: "Death", "Pandemic", "Disease", "Illness"

5.3 Topic features from prior research

Leveraging prior studies, the following features were generated by topic using keyword. Specific keywords used for each topic can be found here:

https://github.com/ajpag/US-News-NLP/blob/53faa3e31d901386b12ea0459586ac6bd0785f1e/analysis/feature_engineering_ap.R#L280

Source: “Politicization and Polarization in COVID-19 News Coverage” <https://journals.sagepub.com/doi/full/10.1177/1075547020950735>

- COVID 19
- Scientist
- Republican
- Democrat

Source: “Polarization in elite communication on the COVID-19 pandemic” <https://advances.sciencemag.org/content/6/28/eabc2717>

- Republican Words
- Democrat Words

5.4 Topic Modeling: Latent Dirichlet Allocation

To assess commonalities in topics and article text across articles, Latend Dirichlet Allocation was applied. After experimentation, it was decided that 7 topics was the optimal number, based on interpretability and overlap of topics. The VEM method was used: <https://www.tidytextmining.com/topicmodeling.html>

Topics:

Table 1: Topics from LDA

topic_number	topic
1	Politics
2	Reported deaths and cases
3	China / Wuhan
4	Outbreaks and infections by country
5	Patients and Symptoms
6	Vaccines and Research
7	Business and Economy

6 Measuring Bias Across News Sources

Performing a chi-square test to measure differences in significances across news sources, the following metrics were used. The chi-square test for all three metrics suggest that there is a indeed a bias in article text across the news sources.

6.1 Topic Probabilities

Following intuition, The Wall Street Journal has the highest average topic probability for Business and Economy articles. It is interesting that Reuters has the highest probability of China / Wuhan related articles. CNN has the lowest probability of articles pertaining to vaccines and research.

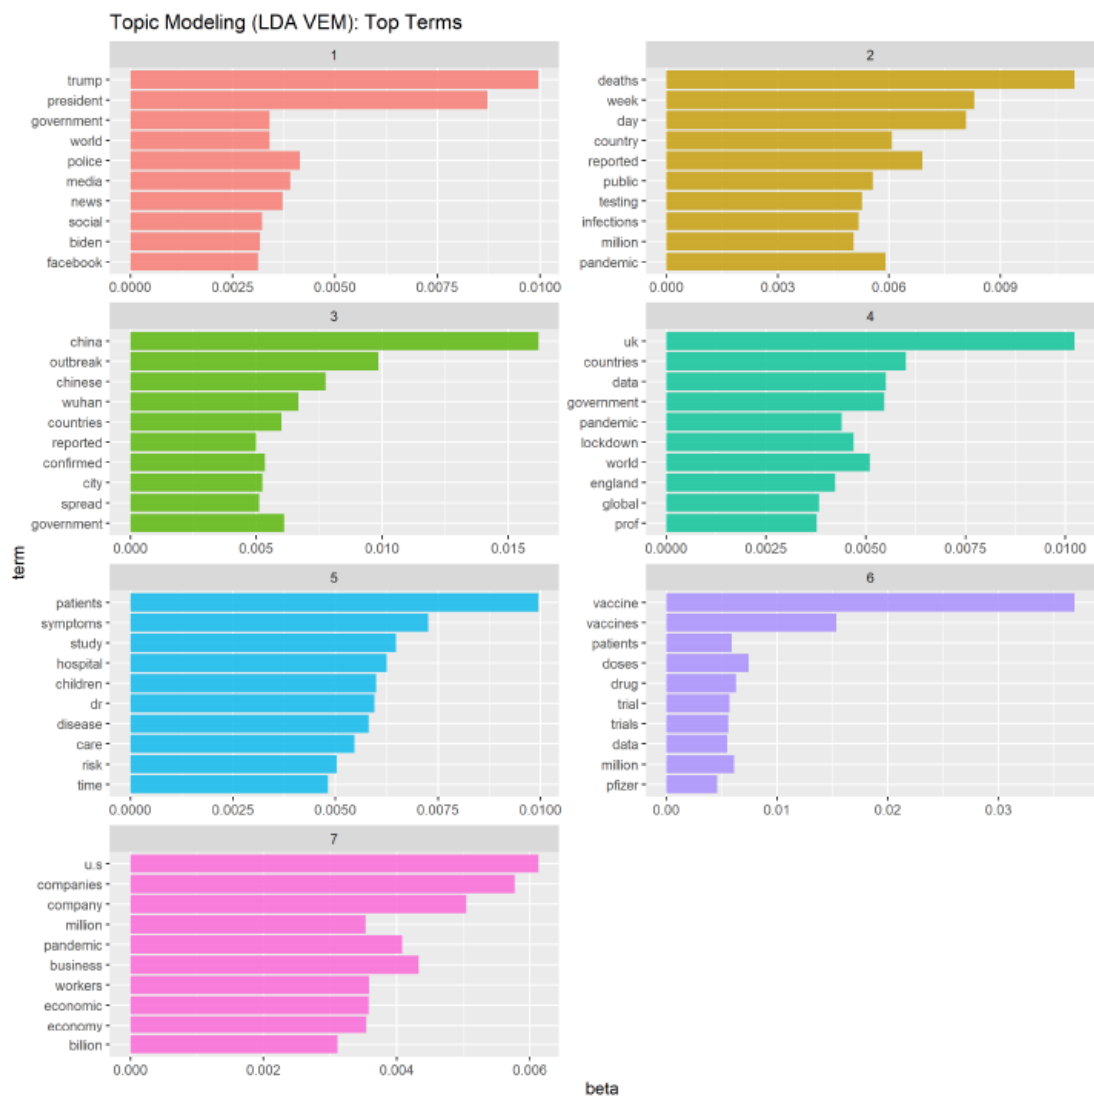


Figure 5: Topic Modeling: Top Words

Figure 5: Topic Modeling: Top Words

The chi-square test ($p < 1\%$) indicates there are correlations across each news source in relation to average topic probabilities.

Note: Probabilities were multiplied by a factor of 100 prior to running the test




Figure 6: Topic Modeling Probabilities

Figure 6: Topic Modeling Probabilities

6.2 Average Sentiment

The chi-square test ($p < 1\%$) indicates there are correlations across each news source in relation to average AFINN sentiment.

Note: Sentiment scores were multiplied by a factor of 100 and converted to positive integers prior to running the test




Figure 7: Avg sentiment by sentence across articles

6.3 Word Count

The chi-square test ($p < 1\%$) indicates there are correlations across each news source in relation to average word count.

Note: Average word counts were converted to positive integers prior to running the test




Figure 8: Avg word count across articles

7 Classification

For each of the 5 algorithms run, 4 models were run using different sets of features:

- **Topic Modeling features:** per section 5.4
- **Keyword features:** per section 5.2
- **Topic keyword features:** per section 5.3
- All features

Table 2: Classification Model Results

model	train_accuracy	test_accuracy	auc
gbm_topic_lda	0.65	0.69	0.89
gbm_keywords	0.64	0.66	0.86
gbm_topic_keywods	0.37	0.43	0.70
gbm_all_features	0.80	0.85	0.97

8 Results

8.1 Conclusion and Next Steps

8.2 Limitations:

- Selection bias in news articles analyzed: Due to legal restrictions, more conservative news sources, such as Fox News, were not scraped.
- Context limitations in sentiment: The sentiment method used, AFINN, only parsed words individually, and not into context of the entire article. So, the sentence sentiment was calculated using the average sentiment of words.
- Parsing limitations: There are many edge cases in which breaking down articles by sentences did not parse successfully. For example, `tidytext::unnest_functions()` incorrectly parsed the following sentence into two since Ms. has a period:

Original sentence: "The manager's decision to send Ms. Coleman home for wearing the headscarf was due to a lack of training," Warren said.

Parsed sentence(s)

- "The manager's decision to send Ms."
- "Coleman home for wearing the headscarf was due to a lack of training," Warren said."